

ADESH KUMAR SAHOO

Kolkata, India | adesh.ks2002@gmail.com
GitHub: github.com/adeshkumarm1 | LinkedIn

SUMMARY

Applied Machine Learning and Natural Language Processing practitioner with hands-on experience in building retrieval-based NLP systems, predictive ML pipelines, and LLM evaluation frameworks. Experienced in modular ML workflows, model evaluation, and real-world deployment.

PROJECTS

LLM and RAG Evaluation Benchmarking Framework

Sentence Transformers, OpenAI API, FastAPI, Streamlit

- Designed and deployed a benchmarking framework for comparing baseline LLM and RAG-based systems for QA evaluation using semantic retrieval and LLM-based generation.
- Built modular retrieval, generation, evaluation, and orchestration pipelines supporting both local and cloud inference workflows.
- Implemented multi-metric benchmarking using Exact Match, Semantic Similarity, Recall@k, MRR, and LLM-as-Judge evaluation.
- Improved semantic similarity by 58% and LLM-judged correctness by 35.8% over baseline responses across a 150-sample benchmark, with 95% Recall@k retrieval performance.

End-to-End Flight Delay Prediction System - Predicting Risk, Severity, and Causes using Pre-departure Information

Scikit-learn, XGBoost, Streamlit, FastAPI

- Designed and deployed an end-to-end flight delay prediction system enabling real-time pre-departure prediction of delay risk and severity through a modular ML pipeline.
- Engineered production-oriented ML inference workflows by resolving training-serving skew, fixing feature-space inconsistencies, and deployment memory constraints, ensuring stable predictions across local and cloud environments.
- Built an aviation-informed multitask prediction system for delay risk, severity, and possible causes using domain-specific delay logic while enforcing strict data-leakage prevention and reusable preprocessing pipelines.
- Evaluated multiple modeling approaches and identified feature richness over model complexity as the primary performance bottleneck, informing future feature-engineering and data expansion efforts.

Context-Aware Business Japanese to English Interpreter Assistant (Ongoing Project)

- Currently building a manually curated bilingual business Japanese-English corpus (~500 sentence pairs) sourced from real business communication and annotated with metadata such as scenario, domain, and formality.
 - Developing a context-aware business communication assistant using multilingual retrieval, metadata-aware ranking, and human-in-the-loop suggestion selection.
-

SKILLS

Languages: Python

Machine Learning: Classification, Regression, Feature Engineering, Model Evaluation, Pipelines

NLP and LLMs: Sentence Transformers, Semantic Retrieval, RAG, Embeddings

Libraries: Scikit-learn, Numpy, Pandas

Deployment: Streamlit, FastAPI, Render

Tools: Git, Jupyter

EDUCATION

Bachelor of Technology in Information Technology

Meghnad Saha Institute of Technology

2020 - 2024 | CGPA: 8.79

CERTIFICATIONS

Deep Learning Specialization - Andrew Ng (Coursera)

Machine Learning Specialization - Andrew Ng (Coursera)

Japanese Language Proficiency Test N2